

Challenge context

Q1-What is the biggest problem in the customer journey today?

Each channel — WhatsApp, contact center, website, and mobile app — operates independently with no shared memory. The consequences are clear: customers repeat themselves every time they switch channels, agents have zero context, resolution is delayed, and leadership has no unified view of what is happening across the ministry. The core failure is fragmentation of the customer journey. The challenge asks you to fix this with a single unified AI brain — not a better chatbot on one channel.

Q2-Where exactly do customers repeat their information?

Repetition happens at every channel transition. The most common pattern in the dataset: a customer submits a request on the website, follows up on WhatsApp to ask for status, then calls the contact center when they hear nothing — and re-explains the full issue each time. Your system must use Customer ID to pull the full unified profile at the start of every interaction on any channel, so the customer never repeats themselves.

Q3-Which channel is used the most today?

The website is the primary channel for this challenge. Most formal service transactions — housing applications, land registration, vehicle permits, boat licensing — start on the website because they require document uploads and structured forms. WhatsApp and the mobile app are important secondary channels, mainly used for status follow-ups, notifications, and quick inquiries. Design your main prototype flow around the website experience first, then show how it connects to the other channels.

Q4-What are the top 5 most-used services?

1) Housing Assistance Requests, 2) Service Request Status Checks, 3) Vehicle & Truck Registration, 4) Boat / Vessel Registration, 5) Land Ownership Transfers. These should be the services your prototype handles end-to-end.

Q5-Which service creates the most complaints and escalations?

Housing Assistance Requests — because customers apply, wait weeks with no update, then call or WhatsApp repeatedly asking for status. Lack of proactive communication is the core issue, not the service itself. Your AI should auto-send status updates to prevent this.

Q6-Which customer type should we focus on first?

The guidelines identify four customer types: individuals, businesses, investors, and government entities. For your prototype, focus on individual residents first — they represent the highest volume of requests and the widest variety of services (housing, vehicle, land, maritime). A solution built for individuals will generalize better in a demo than one built for a niche segment. Businesses and investors are important but lower frequency and typically have dedicated account managers.

Q7-If you can build only one complete journey for the demo, which one?

Build the Housing Maintenance Request journey — it touches every requirement of the challenge in one flow:

Customer visits the website → AI detects intent from the page + search → bot collects details and creates a case → customer gets a confirmation → days later, customer asks for status on WhatsApp → system pulls the unified profile and gives an instant update without asking who they are → if the case is delayed, AI proactively sends a notification → if escalated, the contact center agent sees the full history on screen.

This single journey demonstrates intent detection, unified profile, cross-channel continuity, case automation, proactive engagement, and human handoff — all in one demo.

Q8-Should customers be able to start on Website, continue by phone, and finish on the app?

Yes — this is the single most important requirement in the entire challenge. The guidelines are explicit: a customer who starts on WhatsApp, continues through the contact center, and completes on the mobile app must feel like they are speaking to the same intelligent representative throughout. Customer ID is your linking key across all datasets. At every channel entry point, the system must retrieve the unified profile first, before doing anything else.

Q9-What must be in the unified customer profile?

The official guidelines specify: personal information, previous requests, complaints, interaction history, satisfaction level, and preferred channel. All of these are available in the CRM sheet of the dataset. Add to that: VIP tier, risk flag (repeat escalator / high churn risk), open cases, cross-channel interaction count, and last service requested. The unified profile is the foundation — every other AI capability depends on it being accurate and instantly accessible.

Q10-How should the system identify the same customer across channels?

Use phone number as the primary real-time identifier — it is present across WhatsApp, voice calls, and app login and requires no extra step from the customer. Emirates ID is the gold standard for formal identity verification but creates friction. In the dataset, use Customer ID as the linking key across all five sheets. For the prototype, a phone number → Customer ID lookup is sufficient and realistic.

Q11-What should we automate first?

The guidelines list the full automation sequence: create → categorize → route → track → close. For the prototype, prioritize in this order:

- 1) Intent detection and FAQ answers — highest volume, fastest win
- 2) Automated case creation — eliminates manual data entry
- 3) Status tracking self-service — eliminates most follow-up contacts
- 4) Smart routing — reduces wrong-queue transfers
- 5) Auto-closing simple cases — reserved for fully self-served interactions

Depth on steps 1–3 is more impressive than shallow coverage of all five.

Q12-Which cases must always involve a human?

Always escalate to a human when: emergency or structural safety risk is reported, legal or inheritance disputes arise, anger flag is triggered (from sentiment analysis), VIP or Platinum tier customer is involved, case has breached its SLA without resolution, or critical documents are missing and the application is blocked. In the dataset, you can use the combination of: Anger Flag = Yes, Sentiment = Very Negative, Priority = Critical, or Reopen Count > 1 as your escalation triggers.

Q13-What should trigger an escalation?

Use these signals from the dataset as your escalation logic: Anger Flag = Yes, Sentiment Label = Very Negative, SLA Met = No, Reopen Count > 1, Risk Flag = Repeat Escalator, Priority = Critical, VIP Tier = Gold or Platinum.

Best practice: combine at least two signals before escalating — a single negative sentiment alone is not enough, but negative sentiment + SLA breach is a clear trigger. The guidelines also highlight the importance of predicting escalations before they happen — this is one of the innovation bonus areas.

Q14-What does the human agent need most from the AI co-pilot?

Note: Part C is out of scope for this challenge

The guidelines specify the co-pilot should provide: best response recommendation, best action, relevant knowledge articles, risk level of the case, and escalation probability.

The single highest-value feature: a real-time customer summary card shown the moment an agent receives a contact — who the customer is, what they want, what has already happened, and the recommended next action. This eliminates "please hold while I check your file" and immediately impresses judges who have worked in customer service operations.

Q15-What does leadership want to see in the dashboard?

Note: Part C is out of scope for this challenge

The official guidelines are specific here: request volumes, customer satisfaction, channel performance, most requested services, escalation risks, and workforce productivity. The innovation ideas section goes further — the "National Customer Experience Command Center" concept shows satisfaction heat maps, escalation hotspots, and demand distribution across the ministry in real time. Even a static mockup of this dashboard in your presentation would earn strong marks under the Analytics & Innovation criterion (10% of the score).

Q16-Which KPI matters most to the ministry?

Two KPIs above all others: Customer Satisfaction (CSAT) — the external-facing measure used in government performance benchmarks — and resolution time, which directly drives CSAT. If your system can demonstrate a measurable reduction in resolution time using the dataset (compare average resolution hours for escalated vs. non-escalated cases, or SLA breach rates), that is the most compelling evidence you can show. Even simulated before/after numbers make a strong impression.

Q17-What proactive updates should the AI send?

The guidelines specify four types: status updates, notifications, reminders, and alerts. In practice for MOEI services:

- 1) "Your maintenance request has been assigned — field visit scheduled for Thursday"
- 2) "Your application is on hold — please upload your title deed to proceed"
- 3) "Your boat license has been approved — collect from the Al Rashidiyah service center"
- 4) "We apologize for the delay — a senior agent will contact you within 24 hours"

Always send via the customer's preferred channel (stored in the unified profile). The "Proactive Suggestion Shown / Accepted" columns in the Web & App sheet give you training signal for when these work.

Q18-What data is available for the hackathon?

The guidelines list six data source categories: customer data (CRM + profiles + service history), interaction data (chatbot, WhatsApp, call logs), operational data (KPIs, queue metrics, handling times, agent logs), digital analytics (website + app + search logs), knowledge sources (FAQs, policies, procedures, service catalogues), and customer feedback (surveys, comments, ratings).

All of this is represented in the five-sheet dataset provided: WhatsApp Logs (300 sessions), Voice Logs (400 calls), Website & Mobile Sessions (350 records), CRM Profiles (200 customers), Service Cases (250 records). Link everything using Customer ID and Case ID.

Q19-Is mock data and simulated APIs acceptable, or do you expect real integration?

Fully acceptable — and expected. The dataset provided is synthetic. What judges evaluate is not whether you connected to real ministry systems, but whether your architecture is credible and your reasoning is sound. Label your mock components clearly in the demo and explain what the real integration point would be (e.g. "this would call the MOEI CRM API via secure government gateway"). This shows maturity and earns trust.

Q20-What are the top 3 things to focus on for a strong winning solution?

Based directly on the evaluation criteria (Omnichannel Design 25%, Customer Experience & Unification 20%, Challenge Understanding 20%):

1. Cross-channel memory that actually works — demonstrate one customer moving across website → WhatsApp → contact center without repeating themselves. This is where the highest-weighted criterion (25%) is won or lost.
2. A unified profile that is the source of truth — every channel action should read from and write to the same customer record. Show this explicitly in your demo, not just in a slide.
3. Proactive over reactive — the guidelines are emphatic: the system should anticipate needs, not just respond to them. Show at least one proactive notification or next-best-action suggestion in your demo flow. Teams that only build a reactive chatbot will lose marks on the Customer Experience criterion.

Q21-How can we provide unique value that will truly impress the judges?

The official guidelines name specific ideas that can earn bonus marks. The strongest ones to consider within hackathon time:

- Customer Digital Twin — a dynamic profile that learns behavior, predicts future needs, and proactively suggests the next service before the customer even asks. Use the CRM risk flags and service history as the foundation.
- Emotion-Aware responses — go beyond positive/neutral/negative. Use the anger flag and 3-point sentiment arc from the voice data to automatically change the tone, urgency, and routing of responses. A frustrated customer gets a different experience than a satisfied one.
- Predictive complaint prevention — use the Reopen Count, SLA breach history, and cross-channel contact frequency in the dataset to predict which cases are about to become complaints — and intervene before the customer contacts again.
- Multi-agent ecosystem — instead of one monolithic AI, show a Housing Agent, Maritime Agent, and Transport Agent all coordinated by a Master Agent. This architecture directly mirrors how the ministry is organized and shows real systems thinking.